

GRAPHICS PROGRAMMING I

OPTIMIZATIONS

Ray Tracing| Optimizations



- Increase the intersection speed
 - More performant Hit Algorithms
 - Moller Trumbore (Triangle Intersection)
 - Sphere HitTest (Geometric vs Analytic)
 - ...
- Reduce the number of intersections
 - Acceleration Structures
 - Axis-Aligned Bounding Boxes (AABB) - SlabTest
 - Uniform Grid Algorithm
 - Bounding Volume Hierarchies (BVH)
 - Octree Algorithm
 - Binary Space Partitioning (BSP) Tree Algorithm
 - KD-Tree Algorithm
 - ...
- Reduce the number of rays
- Data vs Object Oriented framework
- Use parallel algorithms
 - Multithreading
 - Thread / ThreadPools
 - Parallel For
 - Async

Ray Tracing| Multithreading

A software implementation allowing different threads to be executed concurrently. A multithreaded program appears to be doing several things at the same time even when it's running on a single-core machine (concurrent).

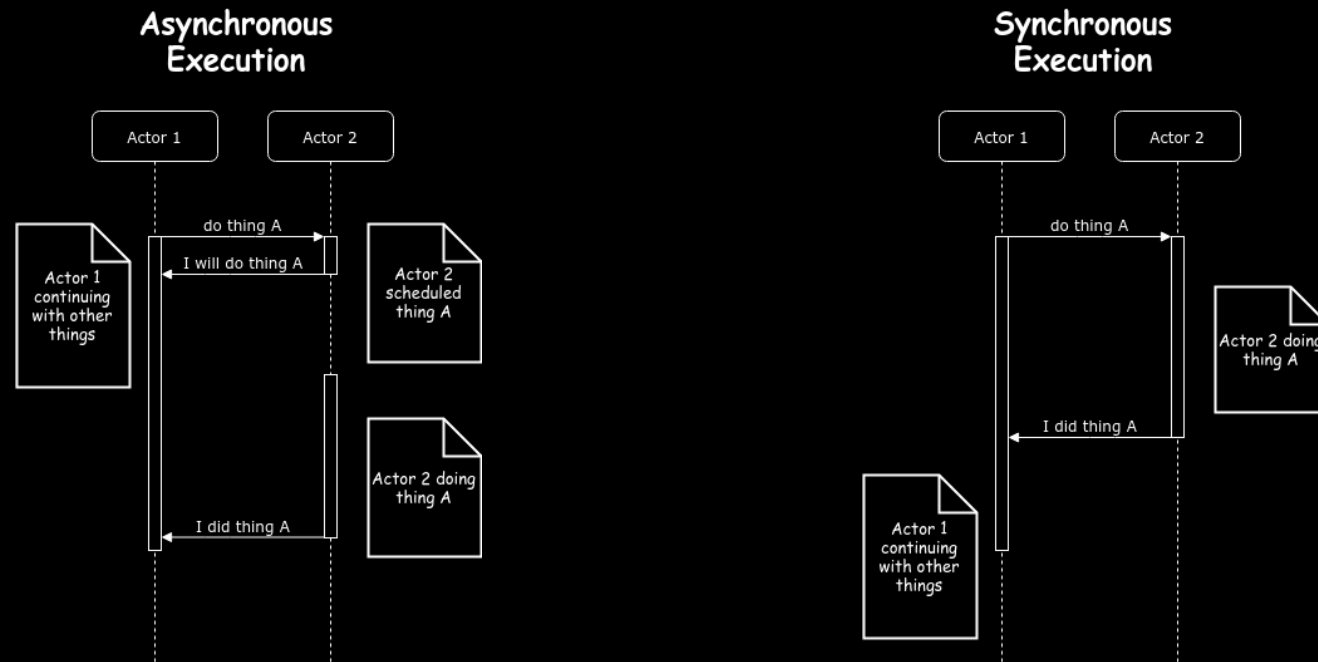
Concurrency vs Parallelism

- **Concurrency**
 - Executing multiple tasks at the same time but not necessarily simultaneously. On a single-core machine it's an illusion of multiple tasks running in parallel because of a very fast switching between threads/tasks by the CPU.
- **Parallelism**
 - When the tasks are actually being executed in parallel on multiple cores.
- Multi-core machines will of course use a combination of concurrency and parallelism



Ray Tracing| Async

- Asynchrony
 - Refers to the fact that one event might be happening at a different time (not in synchrony) to another event.
 - Depending on the execution context, async execution can run in parallel (std::async > Depends on the context, compiler & operating system)
 - Commonly used to execute tasks on a separate thread to prevent blocking the main thread (ex. UI Responsiveness)




Ray Tracing| Hands-On

1. Restructure 'renderer' > Function to render a single pixel (task)
2. Synchronous Execution (default)
3. Asynchronous Execution (std::async)
 1. Limited to the number of logical cores (semi-parallel)
4. Parallel Execution (std::parallel_for)

```
C:\Users\thoma\Desktop\GP1_W3
dFPS: 7.40012
dFPS: 7.46454
dFPS: 7.3773
dFPS: 6.91264
dFPS: 7.47469
dFPS: 7.24606
dFPS: 7.26912
dFPS: 7.65877
dFPS: 7.156
dFPS: 7.38399
dFPS: 7.30292
dFPS: 7.52395
dFPS: 7.63575
dFPS: 7.66931
dFPS: 7.20711
dFPS: 6.79094
dFPS: 7.26584
```

```
C:\Users\thoma\Desktop\GP1_W3
dFPS: 36.6462
dFPS: 37.8521
dFPS: 38.8462
dFPS: 45.2962
dFPS: 43.8583
dFPS: 45.6293
dFPS: 50.6976
dFPS: 50.713
dFPS: 49.9908
dFPS: 51.0918
dFPS: 42.6724
dFPS: 38.8834
dFPS: 38.8616
dFPS: 33.2757
dFPS: 37.9614
```

```
C:\Users\thoma\Desktop\GP1_W3
dFPS: 53.1368
dFPS: 56.3408
dFPS: 54.1007
dFPS: 57.8879
dFPS: 57.9573
dFPS: 56.7208
dFPS: 56.0423
dFPS: 57.2159
dFPS: 55.822
dFPS: 56.6522
dFPS: 57.7766
dFPS: 56.901
```



Ray Tracing| Hands-On (MultiThreading)

- Create function to process a single pixel (Renderer::RenderPixel)

Renderer.h

```
void Render(Scene* pScene) const;

void RenderPixel(Scene* pScene, uint32_t pixelIndex, float fov, float aspectRatio, const Camera& camera, const std::vector<Light>& lights, const std::vector<Material*>& materials) const;

bool SaveBufferToImage() const;
```

Renderer.cpp

```
void Renderer::RenderPixel(Scene* pScene, uint32_t pixelIndex, float fov, float aspectRatio,
    const Camera& camera, const std::vector<Light>& lights, const std::vector<Material*>& materials) const
{
    const int px = pixelIndex % m_Width;
    const int py = pixelIndex / m_Width;

    float rx = px + 0.5f;
    float ry = py + 0.5f;

    float cx = (2 * (rx / float(m_Width)) - 1) * aspectRatio * fov;
    float cy = (1 - (2 * (ry / float(m_Height)))) * fov;

    //Same code you already have but without running it inside a double for-loop
    //...
```

Ray Tracing| Hands-On (MultiThreading)

- `Renderer::Render`
 - The majority of the code can be removed (double for-loop)
 - Data that every pixel needs for processing can be prefetched here
 - Camera (& cameraToWorld matrix calculation)
 - Materials
 - Lights
 - Aspect Ratio
 - Field of View
 - Calculate the total number of pixels (width * height)
 - Make use of 'preprocessor directives' to switch between implementations

Renderer.cpp (top)

```
//#define ASYNC  
//#define PARALLEL_FOR
```

Renderer.cpp

```
void Renderer::Render(Scene* pScene) const  
{  
    Camera& camera = pScene->GetCamera();  
    camera.CalculateCameraToWorld();  
  
    const float fovAngle = camera.fovAngle * TO_RADIANS;  
    const float fov = tan(fovAngle / 2.f);  
  
    const float aspectRatio = m_Width / static_cast<float>(m_Height);  
  
    auto& materials :const vector<Material*>& = pScene->GetMaterials();  
    auto& lights :const vector<Light*>& = pScene->GetLights();  
  
    const uint32_t numPixels = m_Width * m_Height;  
  
    #if defined(ASYNC)  
        //Async Logic  
    #elif defined(PARALLEL_FOR) #if defined(ASYNC)  
        //Parallel-For Logic  
    #else #elif defined(PARALLEL_FOR)  
        //Synchronous Logic (no threading)  
    #endif #elif defined(PARALLEL_FOR) #else  
  
        //END  
        //Update SDL Surface  
        SDL_UpdateWindowSurface(m_pWindow);  
    }  
}
```

Ray Tracing| Hands-On (MultiThreading)

- Synchronous Logic (No Threading)
 - Simple for-loop that loops numPixels amount of times

Renderer::Render

```
#else #elif defined(PARALLEL_FOR)

    //No Threading
    //+++++++
    for (uint32_t i{ 0 }; i < numPixels; ++i)
    {
        RenderPixel(pScene, i, fov, aspectRatio, pScene->GetCamera(), lights, materials);
    }

#endif #elif defined(PARALLEL_FOR) #else
```


Ray Tracing| Hands-On (MultiThreading)

- Async Logic (using futures) [1]

- Similar to the 'synchronous' approach, but now we are spreading the work across multiple async operations
- Systems decides how these are scheduled (could be parallel, could be concurrent) – depends on system resources
- We are limiting the amount of async operations to the total number of cores we have available
- Requires the 'future' header

```
#define ASYNC  
//#define PARALLEL_FOR
```

Renderer.cpp

```
#include <future> //async
```

Renderer::Render (1)

```
#if defined(ASYNC)  
    //Async  
    //+++++  
    const uint32_t numCores = std::thread::hardware_concurrency();  
    std::vector<std::future<void>> async_futures{};  
    const uint32_t numPixelsPerTask = numPixels / numCores;  
    uint32_t numUnassignedPixels = numPixels % numCores;  
    uint32_t currPixelIndex = 0;
```

Ray Tracing| Hands-On (MultiThreading)

- Async Logic (using futures) [2]

Renderer::Render (2)

```
for(uint32_t coreId{0}; coreId < numCores; ++coreId)
{
    uint32_t taskSize = numPixelsPerTask;
    if (numUnassignedPixels > 0)
    {
        ++taskSize;
        --numUnassignedPixels;
    }

    async_futures.push_back(_Val: std::async(_Policy: std::launch::async, _Fnarg: [=, this] ~>void
    {
        //Render all pixels for this task (currPixelIndex > currPixelIndex + taskSize)
        const uint32_t pixelIndexEnd = currPixelIndex + taskSize;
        for (uint32_t pixelIndex{ currPixelIndex }; pixelIndex < pixelIndexEnd; ++pixelIndex)
        {
            RenderPixel(pScene, pixelIndex, fov, aspectRatio, camera, lights, materials);
        }
    }));

    currPixelIndex += taskSize;
}

//Wait for async completion of all tasks
for (const std::future<void>& f : async_futures)
{
    f.wait();
}

#elif defined(PARALLEL_FOR) #if defined(ASYNC)
```

Ray Tracing| Hands-On (MultiThreading)

- Parallel Logic

- We can easily parallelize our for-loop by making use of the “Parallel Patterns Library” (PPL) which is a part of the C++ Standard Library
- Same as with the Async approach, the system is in charge of scheduling the tasks but ensure execution in a parallel manner depending on the available number of computing resources. (Should give you better performance compared to the async approach)
<https://learn.microsoft.com/en-us/cpp/parallel/concrt/parallel-algorithms?view=msvc-170>
- Requires the ‘ppl’ header (+ parallel functions are part of the ‘concurrency’ namespace)

```
//#define ASYNC  
#define PARALLEL_FOR
```

Renderer.cpp

```
#include <ppl.h> //parallel_for
```


Renderer::Render

```
#elif defined(PARALLEL_FOR) #if defined(ASYNC)  
  
    //Parallel For  
    //+++++  
    concurrency::parallel_for(_First:0u, _Last:numPixels, _Func:[=, this](int i)->void {  
        RenderPixel(pScene, i, fov, aspectRatio, camera, lights, materials);  
    });  
  
#else #elif defined(PARALLEL_FOR)
```

- **parallel_for (concurrency namespace)**
 - Works very similar to a traditional for-loop (iterating over a specified range) but with automated parallelization under the hood



Ray Tracing| Acceleration Structure

- An acceleration structure is super useful to prevent unnecessary calculations (= reducing the number of intersection test per frame)
- There are lots of acceleration techniques available
- For this hands-on we are going to implement a very simple one, **The Slab-Test**
 - > AABB bases pre-hit check for our TriangleMeshes
 - > AABB (Axis-Aligned BoundingBox)



Ray Tracing| Acceleration Structure

- Slab Test (AABB-Ray Intersection)
 1. Define a AABB around the TriangleMesh (Min/Max)
 2. Perform a HitTest with the AABB box
 1. Test each 'slab'
 3. If NO HIT > Do not HitTest the TriangleMesh
 4. IF HIT > Perform TriangleMesh HitTest



$$t0x = (B0_x - O_x) / D_x$$

$$t1x = (B1_x - O_x) / D_x$$

$$t0y = (B0_y - O_y) / D_y$$

$$t1y = (B1_y - O_y) / D_y$$

$$t0z = (B0_z - O_z) / D_z$$

$$t1z = (B1_z - O_z) / D_z$$

Ray Tracing| Acceleration Structure

- Update Vector3 class
 - Vector3::Min > Get smallest components of 2 vectors
 - Vector3::Max > Get biggest components of 2 vectors

Vector3.h

```
static Vector3 Reflect(const Vector3& v1, const Vector3& v2);  
  
static Vector3 Max(const Vector3& v1, const Vector3& v2);  
static Vector3 Min(const Vector3& v1, const Vector3& v2);
```

Vector3.cpp

```
Vector3 Vector3::Max(const Vector3& v1, const Vector3& v2)  
{  
    return {  
        _x: std::max(_Left: v1.x, _Right: v2.x),  
        _y: std::max(_Left: v1.y, _Right: v2.y),  
        _z: std::max(_Left: v1.z, _Right: v2.z)  
    };  
}
```

Vector3.cpp

```
Vector3 Vector3::Min(const Vector3& v1, const Vector3& v2)  
{  
    return {  
        _x: std::min(_Left: v1.x, _Right: v2.x),  
        _y: std::min(_Left: v1.y, _Right: v2.y),  
        _z: std::min(_Left: v1.z, _Right: v2.z)  
    };  
}
```

Ray Tracing| Acceleration Structure

- Update TriangleMesh
 - Find AABB for (object-space) vertices (AABB is defined by a Min & Max position)
 - Calculate the AABB after transforming the vertices (world-space)

DataType.h (TriangleMesh)

```
Matrix scaleTransform{};

Vector3 minAABB;
Vector3 maxAABB;

Vector3 transformedMinAABB;
Vector3 transformedMaxAABB;


std::vector<Vector3> transformedPositions{};
```

Ray Tracing| Acceleration Structure

- Update TriangleMesh
 - Function to calculate the object-space AABB (TriangleMesh::UpdateAABB)
 - Function to calculate the world-space AABB (TriangleMesh::UpdateTransformedAABB) (closest fitting AABB after transformations)

DataType.h (TriangleMesh)

```
void UpdateAABB()  
{  
    //Update AABB Logic  
}  
  
void UpdateTransformedAABB(const Matrix& finalTransform)  
{  
    //Update Transformed AABB  
}
```



Ray Tracing| Acceleration Structure

- TriangleMesh::UpdateAABB

DataType.h (TriangleMesh)

```
void UpdateAABB()
{
    if (positions.size() > 0)
    {
        minAABB = positions[0];
        maxAABB = positions[0];
        for (auto& p:Vector3& : positions)
        {
            minAABB = Vector3::Min(p, minAABB);
            maxAABB = Vector3::Max(p, maxAABB);
        }
    }
}
```

Ray Tracing| Acceleration Structure

- TriangleMesh::UpdateAABB
- Make sure to Update the AABB every time the TriangleMesh::UpdateTransforms is called

TriangleMesh::UpdateTransforms

```
for (auto& p:Vector3& : positions)
{
    transformedPositions.emplace_back(
        finalTransform.TransformPoint(p));
}

// Update AABB
UpdateTransformedAABB(finalTransform);
```

DataTypes.h (TriangleMesh)

```
void UpdateTransformedAABB(const Matrix& finalTransform) {
    // AABB update: be careful -> transform the 8 vertices of the aabb
    // and calculate new min and max.
    Vector3 tMinAABB = finalTransform.TransformPoint(minAABB);
    Vector3 tMaxAABB = tMinAABB;
    // (xmax,ymin,zmin)
    Vector3 tAABB = finalTransform.TransformPoint(maxAABB.x, minAABB.y, minAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmax,ymin,zmax)
    tAABB = finalTransform.TransformPoint(maxAABB.x, minAABB.y, maxAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmin,ymin,zmax)
    tAABB = finalTransform.TransformPoint(minAABB.x, minAABB.y, maxAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmin,ymax,zmin)
    tAABB = finalTransform.TransformPoint(minAABB.x, maxAABB.y, minAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmin,ymax,zmax)
    tAABB = finalTransform.TransformPoint(maxAABB.x, maxAABB.y, minAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmax,ymax,zmin)
    tAABB = finalTransform.TransformPoint(maxAABB.x, maxAABB.y, minAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmax,ymax,zmax)
    tAABB = finalTransform.TransformPoint(maxAABB);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);
    // (xmin,ymax,zmax)
    tAABB = finalTransform.TransformPoint(minAABB.x, maxAABB.y, minAABB.z);
    tMinAABB = Vector3::Min(tAABB, tMinAABB);
    tMaxAABB = Vector3::Max(tAABB, tMaxAABB);

    transformedMinAABB = tMinAABB;
    transformedMaxAABB = tMaxAABB;
}
```

Ray Tracing| Acceleration Structure

- Utils::SlabTest_TriangleMesh

Utils.h

```
inline bool SlabTest_TriangleMesh(const TriangleMesh& mesh, const Ray& ray)
{
    float tx1 = (mesh.transformedMinAABB.x - ray.origin.x) / ray.direction.x;
    float tx2 = (mesh.transformedMaxAABB.x - ray.origin.x) / ray.direction.x;

    float tmin = std::min(_Left: tx1, _Right: tx2);
    float tmax = std::max(_Left: tx1, _Right: tx2);

    float ty1 = (mesh.transformedMinAABB.y - ray.origin.y) / ray.direction.y;
    float ty2 = (mesh.transformedMaxAABB.y - ray.origin.y) / ray.direction.y;

    tmin = std::max(_Left: tmin, _Right: std::min(_Left: ty1, _Right: ty2));
    tmax = std::min(_Left: tmax, _Right: std::max(_Left: ty1, _Right: ty2));

    float tz1 = (mesh.transformedMinAABB.z - ray.origin.z) / ray.direction.z;
    float tz2 = (mesh.transformedMaxAABB.z - ray.origin.z) / ray.direction.z;

    tmin = std::max(_Left: tmin, _Right: std::min(_Left: tz1, _Right: tz2));
    tmax = std::min(_Left: tmax, _Right: std::max(_Left: tz1, _Right: tz2));

    return tmax > 0 && tmax >= tmin;
}
```

Ray Tracing| Acceleration Structure

- Perform SlabTest before testing the entire TriangleMesh

Utils.h (HitTest_TriangleMesh)

```
inline bool HitTest_TriangleMesh(const TriangleMesh& mesh, const Ray& ray, HitRecord& hitRecord, bool ignoreHitRecord = false)
{
    // slabtest
    if (!SlabTest_TriangleMesh(mesh, ray)) {
        return false;
    }

    Triangle triangle{};
}
```

Ray Tracing| Acceleration Structure

- Use in your scenes!

Scene_W4_Bunny::Initialize

```
m_pMesh = AddTriangleMesh(TriangleCullMode::BackFaceCulling, matLambert_White);
Utils::ParseOBJ( filename: %s "Resources/lowpoly_bunny2.obj",
    [&] m_pMesh->positions,
    [&] m_pMesh->normals,
    [&] m_pMesh->indices);

m_pMesh->Scale( %s{ _x:2.f, _y:2.f, _z:2. });

m_pMesh->UpdateAABB();
m_pMesh->UpdateTransforms();
```